



HybridFormer

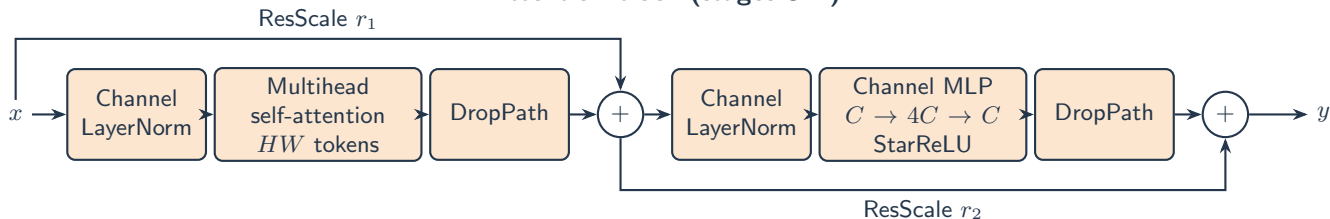
Repository's experimental CIFAR adaptation



Pooling block (stages 1–2)



Attention block (stages 3–4)



Shapes are $H \times W \times C$; s is convolution stride; K is the class count (default 100).

GroupNorm uses one group; channel LayerNorm normalizes C per spatial location.

Attention uses 4 / 6 heads in stages 3 / 4 (64 channels per head).

Channel MLP uses two 1×1 convolutions; StarReLU is $a \text{ReLU}(z)^2 + b$.

ResScale multiplies skip paths and starts at 1; DropPath increases from 0 to 0.1 (training only).