



PoolFormer-S12



PoolFormer block (all stages)



Shapes are $H \times W \times C$; s is convolution stride; K is the class count (default 1000).

GroupNorm uses one group. Channel MLP uses two 1×1 convolutions.

LayerScale starts at 10^{-5} ; DropPath increases from 0 to 0.1 and is disabled at inference.